

Hadoop Deployment

Monday 21st February, 2022

Contents

1	Preinstall	2
1.1	ssh connection	2
1.2	java installation	4
2	Installing Hadoop	4
2.1	hadoop version	4
2.2	download hadoop	4
2.3	extract hadoop tar	5
2.4	move extracted file	6
3	Configuring Hadoop's Java Home	6
3.1	option 1: set a static value	6
3.2	option 2: use readlink to set the value dynamically	6
3.3	set the value on environment	7
4	Create Hadoop User	7
4.1	check your ip address	8
5	On Hadoop User	10
5.1	ssh	11
6	Add Hadoop Services	12
6.1	core-site.xml	12
6.2	hdfs-site.xml	13
6.3	copy to workers	14
7	HDFS	16
7.1	format hdfs	16
7.2	start hdfs	16
8	Resources	17

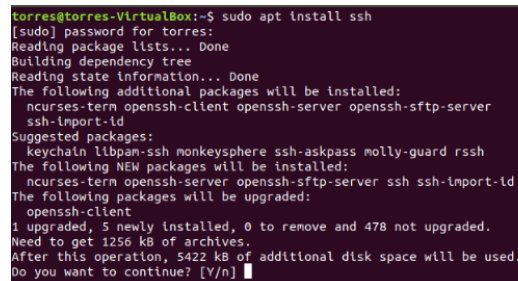
1 Preinstall

1.1 ssh connection

pdsh can run multiple remote commands in parallel.

Listing 1: Install SSH and PDSH

```
sudo apt install ssh
```

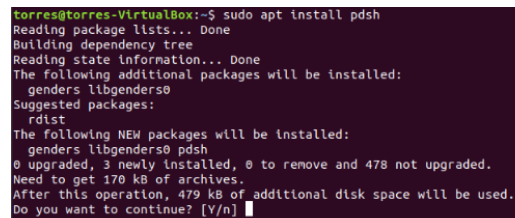


```
torres@torres-VirtualBox:~$ sudo apt install ssh
[sudo] password for torres:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  ncurses-term openssh-client openssh-server openssh-sftp-server
  ssh-import-id
Suggested packages:
  keychain libpam-ssh monkeysphere ssh-askpass molly-guard rssh
The following NEW packages will be installed:
  ncurses-term openssh-server openssh-sftp-server ssh ssh-import-id
The following packages will be upgraded:
  openssh-client
1 upgraded, 5 newly installed, 0 to remove and 478 not upgraded.
Need to get 1256 kB of archives.
After this operation, 5422 kB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Figure 1

Listing 2: Install SSH and PDSH

```
sudo apt install pdsh
```



```
torres@torres-VirtualBox:~$ sudo apt install pdsh
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  genders libgenders0
Suggested packages:
  rdist
The following NEW packages will be installed:
  genders libgenders0 pdsh
0 upgraded, 3 newly installed, 0 to remove and 478 not upgraded.
Need to get 170 kB of archives.
After this operation, 479 kB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Figure 2

Listing 3: Install PDSH

```
nano .bashrc
```

At the end of the file just write the following line:

```
export PDSH_RCMD_TYPE=ssh
```

Now let's configure SSH. Let's create a new key using the following command:

```

torres@torres-VirtualBox: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 .bashrc Modified

alias alert='notify-send --urgency=low -i "${[ $? = 0 ]} && echo terminal || ec$

# Alias definitions.
# You may want to put all your additions into a separate file like
# ~/.bash_aliases, instead of adding them here directly.
# See /usr/share/doc/bash-doc/examples in the bash-doc package.

if [ -f ~/.bash_aliases ]; then
    . ~/.bash_aliases
fi

# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
    if [ -f /usr/share/bash-completion/bash_completion ]; then
        . /usr/share/bash-completion/bash_completion
    elif [ -f /etc/bash_completion ]; then
        . /etc/bash_completion
    fi
fi

export PDSH_RCMD_TYPE=ssh

```

Figure 3

Listing 4: ssh-keygen

```
ssh-keygen -t rsa -P ""
```

```

torres@torres-VirtualBox:~$ ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/home/torres/.ssh/id_rsa):
Your identification has been saved in /home/torres/.ssh/id_rsa.
Your public key has been saved in /home/torres/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:4mA25iJkH2jFujz3ZThFq+41knubh5kvr07ChyA1os torres@torres-VirtualBox
The key's randomart image is:
+---[RSA 2048]---+
|.O      . 0+0   |
|o . = . ..00   |
| . =0+  o .    |
|...000. o =    |
|o . =0. S B =   |
| . o o . o + *  |
| . o E B .      |
| . . + +        |
| . . . +        |
+---[SHA256]---+
torres@torres-VirtualBox:~$

```

Figure 4

Listing 5: ssh-keygen

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

```

torres@torres-VirtualBox:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
torres@torres-VirtualBox:~$

```

Figure 5

Now we can verify the SSH configuration by connecting to the localhost:

Listing 6: ssh local host

```
ssh localhost
```

```

root@hadoop-master:~# ssh localhost
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.4.0-99-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Wed Feb 16 13:54:00 UTC 2022

System load:  0.0               Users logged in:  1
Usage of /:   6.8% of 48.29GB   IPv4 address for eth0: 209.97.140.67
Memory usage: 10%              IPv4 address for eth0: 10.16.0.9
Swap usage:   0%               IPv4 address for eth1: 10.106.0.7
Processes:   111

15 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Last login: Wed Feb 16 13:31:51 2022 from 162.243.188.66
root@hadoop-master:~#

```

Figure 6

1.2 java installation

Install OpenJDK, the default Java Development Kit

Listing 7: jdk

```
sudo apt install default-jdk
```

This step isn't really a step, it's just to check if Java is now correctly installed:

Listing 8: version check

```
java -version
```

```

root@hadoop-master:~# java -version
openjdk version "1.8.0_312"
OpenJDK Runtime Environment (build 1.8.0_312-8u312-b07-0ubuntu1-20.04-b07)
OpenJDK 64-Bit Server VM (build 25.312-b07, mixed mode)
root@hadoop-master:~#

```

Figure 7

2 Installing Hadoop

2.1 hadoop version

In example I used the hadoop-3.0.3 version but we can use any version that we want to change the version in this link:

<https://archive.apache.org/dist/hadoop/core/hadoop-3.0.3/>

2.2 download hadoop

Download Hadoop using the following command:

Listing 9: version check

```
sudo wget https://archive.apache.org/dist/hadoop/core/hadoop-3.0.3/hadoop-3.0.3.tar.gz
```

```
sudo wget
https://archive.apache.org/dist/hadoop/core/hadoop-3.0.3/hadoop-3.0.3.tar.gz.mds
```

Figure 8

```
shasum -a 256 hadoop-3.0.3.tar.gz
```

```
cat hadoop-3.0.3.tar.gz.mds
```

Figure 9

Now that we've verified that the file wasn't corrupted or changed, we'll use the `tar` command with the `-x` flag to extract, `-z` to uncompress, `-v` for verbose output, and `-f` to specify that we're extracting from a file. Use tab-completion or substitute the correct version number in the command below:

```
tar -xzf hadoop-3.0.3.tar.gz
```

Figure 10

5

2.4 move extracted file

Listing 14: Move

```
sudo mv hadoop-3.0.3 /usr/local/hadoop
```

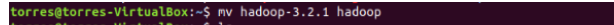


Figure 11

3 Configuring Hadoop's Java Home

The path to Java, `/usr/bin/java` is a symlink to `/etc/alternatives/java`, which is in turn a symlink to default Java binary. We will use **readlink** with the `-f` flag to follow every symlink in every part of the path, recursively. Then, we'll use `sed` to trim `bin/java` from the output to give us the correct value for JAVA HOME.

To find the default Java path:

Listing 15: java path

```
readlink -f /usr/bin/java | sed "s:bin/java::"
```

We can copy this output to set Hadoop's Java home to this specific version, which ensures that if the default Java changes, this value will not.

To begin, open **hadoop-env.sh**:

Listing 16: java path

```
sudo nano /usr/local/hadoop/etc/hadoop/hadoop-env.sh
```

3.1 option 1: set a static value

Listing 17: /usr/local/hadoop/etc/hadoop/hadoop-env.sh

```
. . .
#export JAVA_HOME=${JAVA_HOME}
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64/
. . .
```

3.2 option 2: use readlink to set the value dynamically

It is preferred setting dynamically.

Listing 18: /usr/local/hadoop/etc/hadoop/hadoop-env.sh

```
. . .
#export JAVA_HOME=${JAVA_HOME}
export JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:bin/java::")
. . .
```

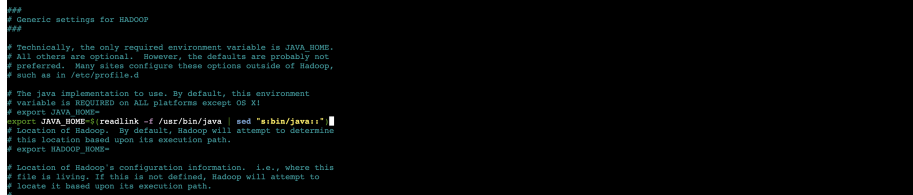


Figure 12

3.3 set the value on environment

Open the environment file on nano with this command:

Listing 19: s

```
sudo nano /etc/environment
```

Then, add the following configurations: (**Path is in one line**)

Listing 20: sudo nano /etc/environment

```
PATH=
"/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/
local/games:/usr/local/hadoop/bin:/usr/local/hadoop/sbin"
JAVA_HOME=$(readlink -f /usr/bin/java | sed "s:bin/java::")
```

4 Create Hadoop User

we will add a user called **hadoopuser**, and we will set up it's configurations:

Listing 21: sudo nano /etc/environment

```
sudo adduser hadoopuser
```

Provide the password and you can leave the rest blank, just press Enter.

```

torres@torres-VirtualBox:~$ sudo adduser hadoopuser
Adding user 'hadoopuser' ...
Adding new group 'hadoopuser' (1001) ...
Adding new user 'hadoopuser' (1001) with group 'hadoopuser' ...
Creating home directory '/home/hadoopuser' ...
Copying files from '/etc/skel' ...
Enter new UNIX password:
Retype new UNIX password:
No password supplied
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
Changing the user information for hadoopuser
Enter the new value, or press ENTER for the default
    Full Name []:
    Room Number []:
    Work Phone []:
    Home Phone []:
    Other []:
Is the information correct? [Y/n] y
torres@torres-VirtualBox:~$

```

Figure 13

Now type these commands :

Listing 22: a

```

sudo usermod -aG hadoopuser hadoopuser
sudo chown hadoopuser:root -R /usr/local/hadoop/
sudo chmod g+rxw -R /usr/local/hadoop/
sudo adduser hadoopuser sudo

```

```

torres@torres-VirtualBox:~$ sudo usermod -aG hadoopuser hadoopuser
torres@torres-VirtualBox:~$ sudo chown hadoopuser:root -R /usr/local/hadoop
torres@torres-VirtualBox:~$ sudo chmod g+rxw -R /usr/local/hadoop
torres@torres-VirtualBox:~$ sudo adduser hadoopuser sudo
Adding user 'hadoopuser' to group 'sudo' ...
Adding user hadoopuser to group sudo
Done.
torres@torres-VirtualBox:~$

```

Figure 14

4.1 check your ip address

Listing 23: a

```
ip addr
```

```

torres@torres-VirtualBox:~$ ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group defau
lt qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP g
roup default qlen 1000
    link/ether 08:00:27:77:85:aa brd ff:ff:ff:ff:ff:ff
    inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic noprefixroute enp0s3
        valid_lft 83598sec preferred_lft 83598sec
    inet6 fe80::ec48:ffb7:4441:13ac/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP g
roup default qlen 1000
    link/ether 08:00:27:e3:97:9c brd ff:ff:ff:ff:ff:ff
    inet 192.168.205.7/24 brd 192.168.205.255 scope global dynamic noprefixroute
enp0s8
        valid_lft 1092sec preferred_lft 1092sec
    inet6 fe80::9fab:96ac:1084:2aff/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
torres@torres-VirtualBox:~$

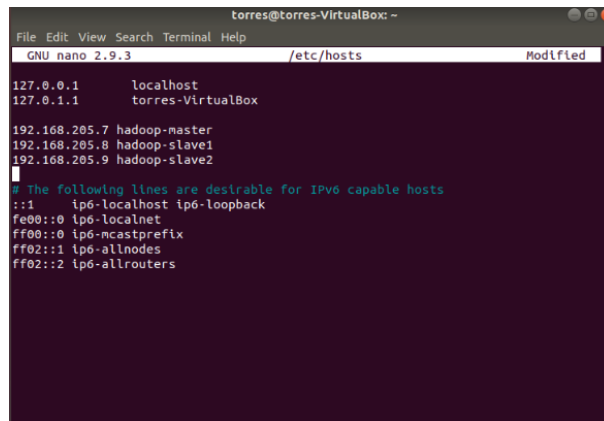
```

Figure 15

Open the hosts file and insert your Network configurations:

Listing 24: a

```
sudo nano /etc/hosts
```

A terminal window titled 'torres@torres-VirtualBox: ~' showing the nano editor editing the file '/etc/hosts'. The file content is as follows:

```
127.0.0.1    localhost
127.0.1.1    torres-VirtualBox

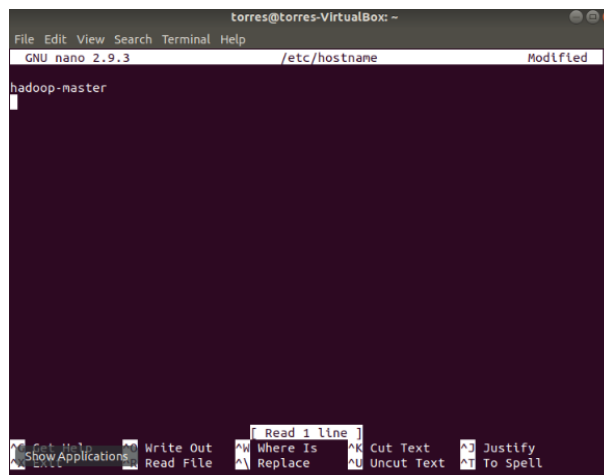
192.168.205.7 hadoop-master
192.168.205.8 hadoop-slave1
192.168.205.9 hadoop-slave2
# The following lines are desirable for IPv6 capable hosts
::1         ip6-localhost ip6-loopback
fe00::0     ip6-localnet
ff00::0     ip6-mcastprefix
ff02::1     ip6-allnodes
ff02::2     ip6-allrouters
```

Figure 16

On the master , open the hostname file on nano:

Listing 25: a

```
sudo nano /etc/hostname
```

A terminal window titled 'torres@torres-VirtualBox: ~' showing the nano editor editing the file '/etc/hostname'. The file content is as follows:

```
hadoop-master
```

Figure 17

You should do the :same for slave nodes:

:




Figure 18




Figure 19

Also, you should reboot all of them so this configuration taked effect:

Listing 26: a

```
sudo reboot
```

5 On Hadoop User

Configure the SSH on hadoop-master, with the hadoopuser. This is the command:

Listing 27: a

```
su - hadoopuser
```




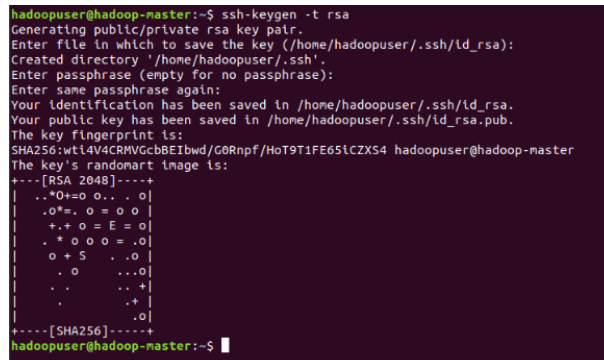
Figure 20

5.1 ssh

Create an SSH key:

Listing 28: a

```
ssh-keygen -t rsa
```



```
hadoopuser@hadoop-master:~$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/hadoopuser/.ssh/id_rsa):
Created directory '/home/hadoopuser/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/hadoopuser/.ssh/id_rsa.
Your public key has been saved in /home/hadoopuser/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:wtI4V4CRMVGcbEibwd/G0Rnpf/Hot9T1FE651CZX54 hadoopuser@hadoop-master
The key's randomart image is:
+---[RSA 2048]---+
| ..O+=0 O.. . O|
| .o*=. o = o o |
| +. + o = E = o|
| . * o o o = .o|
| o + S . .o |
| . o . .o |
| . . . .+ |
| . . . .+ |
| . . . .+ |
| . . . .+ |
+---[SHA256]-----+
hadoopuser@hadoop-master:~$
```

Figure 21

Now we need to copy the SSH key to all the users. Use this command:

Listing 29: a

```
ssh-copy-id hadoopuser@hadoop-master
```



```
hadoopuser@hadoop-master:~$ ssh-copy-id hadoopuser@hadoop-master
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hadoopuser/.ssh/id_rsa.pub"
The authenticity of host 'hadoop-master (192.168.205.7)' can't be established.
ECDSA key fingerprint is SHA256:Nsjcx3SrmwVnSNcXvLYIajjHRdwaEI+RGelkVyoHI4.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- If you are promp
ted now it is to install the new keys
hadoopuser@hadoop-master's password:

Number of key(s) added: 1

Now try logging into the machine, with:  'ssh 'hadoopuser@hadoop-master''
and check to make sure that only the key(s) you wanted were added.
hadoopuser@hadoop-master:~$
```

Figure 22

Listing 30: a

```
ssh-copy-id hadoopuser@hadoop-slave1
```

Listing 31: a

```
ssh-copy-id hadoopuser@hadoop-slave2
```

```

hadoopuser@hadoop-master:~$ ssh-copy-id hadoopuser@hadoop-slave1
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hadoopuser
/.ssh/id_rsa.pub"
The authenticity of host 'hadoop-slave1 (192.168.205.8)' can't be established.
ECDSA key fingerprint is SHA256:Nsjcx3SrmwVnSNcWxvLYiajjHRdwaET+RGelkVyoHI4.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- If you are promp
ted now it is to install the new keys
hadoopuser@hadoop-slave1's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'hadoopuser@hadoop-slave1'"
and check to make sure that only the key(s) you wanted were added.

hadoopuser@hadoop-master:~$

```

Figure 23

```

hadoopuser@hadoop-master:~$ ssh-copy-id hadoopuser@hadoop-slave2
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hadoopuser
/.ssh/id_rsa.pub"
The authenticity of host 'hadoop-slave2 (192.168.205.9)' can't be established.
ECDSA key fingerprint is SHA256:Nsjcx3SrmwVnSNcWxvLYiajjHRdwaET+RGelkVyoHI4.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter
out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- If you are promp
ted now it is to install the new keys
hadoopuser@hadoop-slave2's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'hadoopuser@hadoop-slave2'"
and check to make sure that only the key(s) you wanted were added.

hadoopuser@hadoop-master:~$

```

Figure 24

6 Add Hadoop Services

6.1 core-site.xml

On hadoop-master, open core-site.xml file on nano:

Listing 32: a

```
sudo nano /usr/local/hadoop/etc/hadoop/core-site.xml
```

```

hadoopuser@hadoop-master:~$ sudo nano /usr/local/hadoop/etc/hadoop/core-site.xm
l
[sudo] password for hadoopuser:

```

Figure 25

Then add the following configurations:

Listing 33: a

```

<configuration>
<property>
<name>fs.defaultFS</name>
<value>hdfs://hadoop-master:9000</value>
</property>
</configuration>

```

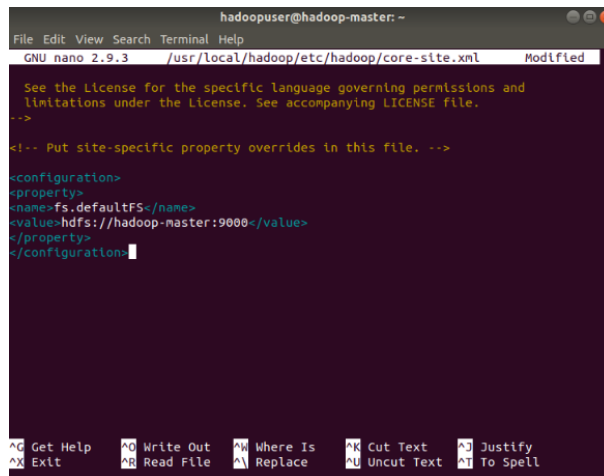


Figure 26

6.2 hdfs-site.xml

Still on hadoop-master, open the hdfs-site.xml file.

Listing 34: a

```
sudo nano /usr/local/hadoop/etc/hadoop/hdfs-site.xml
```


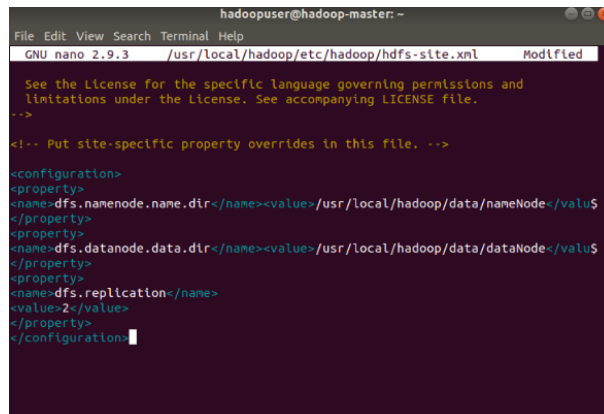


Figure 27

Add the following configurations:

Listing 35: a

```
<configuration>
<property>
<name>dfs.namenode.name.dir</name><value>/usr/local/hadoop/data/nameNode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name><value>/usr/local/hadoop/data/dataNode</value>
</property>
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
</configuration>
```



```
hadoopuser@hadoop-master: ~  
File Edit View Search Terminal Help  
GNU nano 2.9.3 /usr/local/hadoop/etc/hadoop/hdfs-site.xml Modified  
  
See the License for the specific language governing permissions and  
limitations under the License. See accompanying LICENSE file.  
-->  
  
<!-- Put site-specific property overrides in this file. -->  
  
<configuration>  
<property>  
<name>dfs.namenode.name.dir</name><value>/usr/local/hadoop/data/naneNode</valu$  
</property>  
<property>  
<name>dfs.datanode.data.dir</name><value>/usr/local/hadoop/data/dataNode</valu$  
</property>  
<property>  
<name>dfs.replication</name>  
<value>2</value>  
</property>  
</configuration>
```

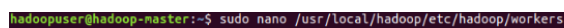
Figure 28

6.3 copy to workers

We're still on hadoop-master, let's open the workers file:

Listing 36: a

```
sudo nano /usr/local/hadoop/etc/hadoop/workers
```



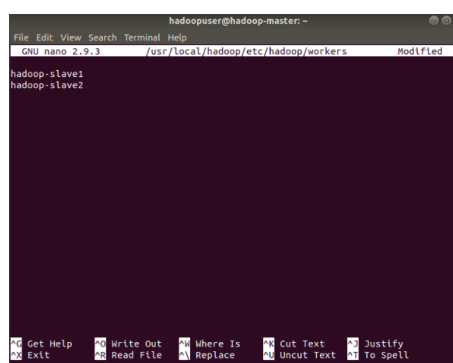
```
hadoopuser@hadoop-master:~$ sudo nano /usr/local/hadoop/etc/hadoop/workers
```

Figure 29

Add these two lines: (the slave names, remember the hosts file?)

Listing 37: a

```
hadoop-slave1  
hadoop-slave2
```



```
hadoopuser@hadoop-master: ~  
File Edit View Search Terminal Help  
GNU nano 2.9.3 /usr/local/hadoop/etc/hadoop/workers Modified  
  
hadoop-slave1  
hadoop-slave2  
  
Get Help Write Out Where Is Cut Text Justify  
Exit Read File Replace Uncut Text To Spell
```

Figure 30

We need to copy the Hadoop Master configurations to the slaves, to do that we use these commands:

Listing 38: a

```
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave1:/usr/local/hadoop/etc/hadoop//
```

```
hadoopuser@hadoop-master:~$ scp /usr/local/hadoop/etc/hadoop/* hadoop-slave1:/u
sr/local/hadoop/etc/hadoop
capacity-scheduler.xml      100% 8260 293.2KB/s  00:00
configuration.xml          100% 1335 791.3KB/s  00:00
container-executor.cfg     100% 1940 1.1MB/s    00:00
core-site.xml              100% 864 776.1KB/s  00:00
hadoop-env.cmd             100% 3999 1.9MB/s    00:00
hadoop-env.sh              100% 16KB 2.6MB/s    00:00
hadoop-metrics2.properties 100% 3321 1.4MB/s    00:00
hadoop-policy.xml          100% 11KB 293.3KB/s  00:00
hadoop-user-functions.sh.example 100% 3414 1.2MB/s    00:00
hdfs-site.xml              100% 1051 201.6KB/s  00:00
https-env.sh               100% 1484 71.0KB/s   00:00
https-log4j.properties     100% 1657 975.4KB/s  00:00
https-signature.secret     100% 21 12.7KB/s   00:00
https-site.xml             100% 620 351.4KB/s  00:00
kms-acls.xml               100% 3518 1.8MB/s    00:00
kms-env.sh                 100% 1351 850.0KB/s  00:00
kms-log4j.properties       100% 1860 990.0KB/s  00:00
kms-site.xml               100% 682 429.3KB/s  00:00
log4j.properties           100% 13KB 11.7MB/s   00:00
mapred-env.cmd             100% 951 646.1KB/s  00:00
mapred-env.sh              100% 1764 1.2MB/s    00:00
mapred-queues.xml.template 100% 4113 2.6MB/s    00:00
mapred-site.xml            100% 758 817.1KB/s  00:00
/usr/local/hadoop/etc/hadoop/shellprofile.d: not a regular file
ssl-client.xml.example     100% 2316 57.9KB/s   00:00
```

Figure 31

Listing 39: a

```
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave2:/usr/local/hadoop/etc/hadoop/
```

```
hadoopuser@hadoop-master:~$ scp /usr/local/hadoop/etc/hadoop/* hadoop-slave2:/u
sr/local/hadoop/etc/hadoop/
capacity-scheduler.xml      100% 8260 1.7MB/s    00:00
configuration.xml          100% 1335 128.3KB/s  00:00
container-executor.cfg     100% 1940 1.9MB/s    00:00
core-site.xml              100% 864 737.1KB/s  00:00
hadoop-env.cmd             100% 3999 3.2MB/s    00:00
hadoop-env.sh              100% 16KB 1.3MB/s    00:00
hadoop-metrics2.properties 100% 3321 2.9MB/s    00:00
hadoop-policy.xml          100% 11KB 1.4MB/s    00:00
hadoop-user-functions.sh.example 100% 3414 2.7MB/s    00:00
hdfs-site.xml              100% 1051 732.5KB/s  00:00
https-env.sh               100% 1484 1.1MB/s    00:00
https-log4j.properties     100% 1657 1.5MB/s    00:00
https-signature.secret     100% 21 15.4KB/s   00:00
https-site.xml             100% 620 529.2KB/s  00:00
kms-acls.xml               100% 3518 2.7MB/s    00:00
kms-env.sh                 100% 1351 987.0KB/s  00:00
kms-log4j.properties       100% 1860 1.8MB/s    00:00
kms-site.xml               100% 682 302.6KB/s  00:00
log4j.properties           100% 13KB 11.4MB/s   00:00
mapred-env.cmd             100% 951 902.9KB/s  00:00
mapred-env.sh              100% 1764 1.6MB/s    00:00
mapred-queues.xml.template 100% 4113 3.8MB/s    00:00
mapred-site.xml            100% 758 799.8KB/s  00:00
/usr/local/hadoop/etc/hadoop/shellprofile.d: not a regular file
ssl-client.xml.example     100% 2316 228.1KB/s  00:00
```

Figure 32

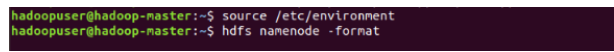
7 HDFS

7.1 format hdfs

Now we need to format the HDFS file system. Run these commands:

Listing 40: a

```
source /etc/environment
hdfs namenode -format
```



```
hadoopuser@hadoop-master:~$ source /etc/environment
hadoopuser@hadoop-master:~$ hdfs namenode -format
```

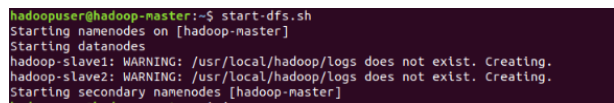
Figure 33

7.2 start hdfs

Start HDFS with this command:

Listing 41: a

```
start-dfs.sh
```




```
hadoopuser@hadoop-master:~$ start-dfs.sh
Starting namenodes on [hadoop-master]
Starting datanodes
hadoop-slave1: WARNING: /usr/local/hadoop/logs does not exist. Creating.
hadoop-slave2: WARNING: /usr/local/hadoop/logs does not exist. Creating.
Starting secondary namenodes [hadoop-master]
```

Figure 34

To check if this worked, run the following command. This will tell you what resources have been initialized:


Listing 42: for all nodes

```
jps
```



```
hadoopuser@hadoop-master:~$ jps
4138 Jps
3771 NameNode
4014 SecondaryNameNode
hadoopuser@hadoop-master:~$
```

Figure 35



```
hadoopuser@hadoop-slave1:~$ jps
1808 DataNode
2024 Jps
hadoopuser@hadoop-slave1:~$
```

Figure 36


```
hadoopuser@hadoop-slave2:~$ jps
1814 DataNode
2031 Jps
hadoopuser@hadoop-slave2:~$
```

Figure 37

Let's see if this worked: Open your browser and type **hadoop-master:9870** This is what mine shows, hopefully yours is showing the same thing!

NomeNode Information - Mozilla Firefox

hadoop-master:9870/dfshealth.html#tab=...

In operation

Show 25 entries

Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Ve
✓ hadoop-slave1:9866 (192.168.205.8:9866)	http://hadoop-slave1:9864	1s	9m	19.56 GB	0	24 KB (0%)	
✓ hadoop-slave2:9866 (192.168.205.9:9866)	http://hadoop-slave2:9864	2s	9m	19.56 GB	0	24 KB (0%)	

Showing 1 to 2 of 2 entries

Previous 1 Next

Figure 38

As you can see, both nodes are operational!

8 Resources

https://medium.com/@jootorres_11979/how-to-set-up-a-hadoop-3-2-1-multi-node-cluster-on-ubuntu-18
<https://www.digitalocean.com/community/tutorials/how-to-install-hadoop-in-stand-alone-mode-on>
<https://archive.apache.org/dist/hadoop/core/hadoop-3.0.3/>
<https://data-flair.training/blogs/install-hadoop-on-ubuntu/>